

What Are Evals and Why We Need Them for Rosetta

A primer on the discipline that separates frontier-grade AI from enterprise embarrassment — and the path forward for Salesforce Data 360

Keenan Vision Working Paper #4

By Vernon Keenan

Spring 2026 · DRAFT v0.1

Who this is for

This book is written for three audiences.

The first is the Salesforce executive who has heard the word “eval” used in product reviews, board presentations, and competitive intelligence briefings, and has been quietly unsure what it means. You are not alone. Most senior leaders in enterprise software in Spring 2026 are in the same position. The vocabulary moved faster than the cultural transmission, and the gap between “we have benchmarks” and “we have an eval culture” is the gap that determines whether your AI products survive the next eighteen months. This book is your shortcut to understanding what your competitors at Anthropic, OpenAI, and Google have known for two years.

The second is the Salesforce architect, admin, or product manager who has been told to “make the agent work better” without being given the instrumentation to know whether it is working. You are doing genuine engineering against an invisible target. This book gives you the vocabulary to name the missing infrastructure and the framework to argue for it.

The third is the partner, customer, or analyst trying to evaluate whether the agent platforms they are buying actually do what the demos suggested. You are the audience the eval economy was built for, even if no one has told you yet.



How to read this book

The full manuscript is roughly forty-five minutes of focused reading. We have organized it so that two reading paths work.

The **executive path** is Chapters 1, 6, and 8. That gets you the definition, the honest diagnosis of where Salesforce sits in Spring 2026, and the action plan. Roughly fifteen minutes. Read these and you will understand what the term means, why the absence of eval culture is a structural problem, and what to do about it next week.

The **practitioner path** is the full sequence. Chapter 2 traces the eighteen-month emergence of the eval economy. Chapter 3 is the technical anatomy. Chapter 4 is the working taxonomy. Chapter 5 is how the frontier labs actually operationalize eval-driven development. Chapter 7 connects the discipline to the Rosetta program specifically. Read all eight in order and you will be able to lead the conversation in your organization rather than be led through it.

Take whichever path your role requires. Do not skip Chapter 6.



Chapter 1 — The Word You Keep Hearing

The word is “eval.” You have heard it in product roadmap reviews, in the hallway after analyst calls, in the LinkedIn posts of competitors you respect, in the questions your most thoughtful customers ask in pre-sales calls. You have probably nodded along once or twice without being entirely sure what was meant. This chapter is the small, specific kindness of giving you a definition you can actually use.

The first thing to understand is that “eval” is currently used in four ways, and the gap between the loosest usage and the strictest is the entire subject of this book. The four uses, ranked from the most casual to the most rigorous, are:

1. The casual demo eval. *“We evaluated the new model and it seems good.”* This is the equivalent of saying you “tested” the brakes on your car by driving down a quiet street and not crashing. It is informal, unrepeatable, and provides no signal you can act on. When a vendor uses “eval” this way in a sales deck, you should treat the claim as having approximately the same epistemic weight as an unverified anecdote.

2. The structured benchmark eval. *“We ran the model against MMLU and it scored 87.3%.”* This is a measurement against a published, fixed test set, with a consistent scoring methodology. Vastly better than the casual version. Still limited, because the test set is public, the model may have been trained on it (a phenomenon called *contamination*), and the benchmark may not reflect your actual use case. This is what most enterprise software vendors mean when they say “eval” today.

3. The preference eval. *“We collected 1,000 human preference judgments comparing the model’s responses against alternatives, with calibrated rubrics and inter-rater reliability above 0.7.”* This is the gold standard for training and improving a model on subjective tasks where there is no single right answer. It is what Mercur’s \$10 billion business sells. It is also what almost no enterprise software vendor in Spring 2026 actually does.

4. The eval-gated release. *“The model failed our internal preference threshold on customer-facing scenarios, so we delayed the release until we could fix it.”* This is eval-driven product development. It requires not just the measurement infrastructure but the institutional authority to block shipment based on what the measurement says. This is what Anthropic, OpenAI, and

Google DeepMind actually do. It is the difference between research and production discipline. It is, almost without exception, what enterprise SaaS companies do not have.

When you hear the word “eval” in the wild, listen for which of these four meanings is operative. The same word covers a 100x range of methodological seriousness, and most of the confusion in current AI conversations comes from people using level 1 or 2 while sounding like they mean level 3 or 4.

The working definition for this book

For the rest of this book, when we use the word “eval,” we mean the following:

An eval is a structured measurement of model behavior against a defined rubric, with verifiable scoring, used to drive product decisions.

Every clause in that sentence carries weight. *Structured* means it is repeatable; you can run it again next week and get a comparable result. *Defined rubric* means there is a written specification of what counts as correct, complete, or preferable; the rubric is the unit of disagreement, not the score. *Verifiable scoring* means the result can be checked by someone other than the original grader; the score is an artifact, not an opinion. *Used to drive product decisions* means there is institutional authority somewhere in the organization to act on the result; without that authority, the eval is research output, not engineering control.

If any of those four elements is missing, you do not have an eval. You have something adjacent to one.

What an eval is not

A surprising amount of confusion dissolves once you can name what evals are *not*.

An eval is not **testing**. Testing checks whether the code does what it was specified to do. Eval measures whether the system produces outcomes the user actually wants. A test passes when the API returns a 200 status code; an eval passes when the API returns the *right* 200 status code, judged against a rubric the testing framework cannot encode.

An eval is not **benchmarking**. Benchmarking is a special case of eval — a frozen, public, contamination-vulnerable special case. All benchmarks are evals; not all evals are benchmarks. You can have a thousand internal evals against your own customer scenarios without ever publishing one as a benchmark, and that is often the right posture.

An eval is not **monitoring**. Monitoring tells you that the model is responding within latency budget, that error rates are within tolerance, that the cost-per-call is tracking forecast. It is

operational instrumentation. It tells you the model is working in the mechanical sense. It does not tell you the model is producing useful answers in the human sense. A model that is hallucinating perfectly within latency budget is a monitoring success and an eval failure.

An eval is not **observability**. Observability lets you trace what happened during a single inference — what was retrieved, what context was assembled, which tools were called. It is forensic infrastructure. It is essential for diagnosing why an eval failed. It is not itself the eval.

An eval is not **QA**. QA, in the traditional enterprise software sense, is a downstream gate run by a separate organization after development claims to be done. The QA team validates against a written specification. Eval, in the AI sense, is upstream and continuous, run by the development team itself, against a rubric that evolves as the product learns what users actually want. Trying to retrofit existing QA functions into an eval discipline is one of the most common and most expensive mistakes enterprise organizations make. The skill set is different, the cadence is different, the authority is different, and the political dynamics are different. Eval is not QA-done-by-AI-people. It is a different practice.

A worked example

The abstract definition above becomes concrete when you walk through a single eval end-to-end. Here is one from the Rosetta program.

The scenario: a Salesforce architect working in Data 360 needs to understand how the Customer Profile entity relates to the Unified Individual record across two acquired company schemas during a post-merger data unification. They ask the agent a question phrased in their own working vocabulary: *“Show me how the customer master from Acme connects to our existing UID structure and flag any field-level conflicts.”*

The eval consists of:

- **The input prompt:** the question above, plus the relevant Data 360 schema context the agent should have access to.
- **Two candidate responses:** Response A, generated by Rosetta in its current production configuration. Response B, generated by Rosetta with a proposed retrieval improvement.
- **The rubric:** a written specification with five criteria. Did the response correctly identify the join keys? Did it surface the field-level conflicts? Did it cite the source schemas with provenance? Did it use the customer’s working vocabulary or default to Salesforce internal jargon? Did it flag governance considerations the customer did not ask about but should know?

- **The grader:** a D360-experienced solution architect, scoring each response against the rubric, expressing a preference (A, B, or tie), and writing a two-to-three-sentence justification explaining the preference.
- **The verifier:** an automated check that confirms the response cited only entities that actually exist in the customer's schema, catching hallucinated field names before the human grader spends time on them.
- **The score:** an aggregate across 1,000 such pairs, broken down by scenario type, retrieval configuration, and rubric criterion, with inter-rater reliability metrics on the 10% of pairs reviewed by a second architect.
- **The decision:** if Response B wins the preference comparison by a defined margin with statistical significance, the retrieval improvement ships. If not, it does not.

That is one eval. It costs roughly \$300 to \$600 in expert time when scaled across the program. Multiply by 1,000 and you have the Phase 2 EvalOps engagement.




Figure 1.1 — Anatomy of an eval, with the seven elements illustrated against the Rosetta D360 worked example. System-generated artifacts in blue, human-authored judgment instruments in amber, the verifier rendered as a hybrid since the worked example uses both LLM-judge and human-expert verification.

Why the vocabulary matters

When leaders use the word “eval” loosely, they think they have a thing they do not have. They believe their engineering team is doing eval-driven development when in fact the team is running occasional benchmark scores in the loosest sense and shipping based on subjective product judgment. The gap between believed eval maturity and actual eval maturity is, in our experience working with major enterprise platforms, almost always at least one full level on the four-level scale above.

This is not a hypothetical concern. Salesforce’s own AI Research team, in October 2025, published the SCUBA benchmark results showing that even closed-source frontier models

complete only 39 percent of realistic CRM tasks in the zero-shot setting. There is no public evidence that this number gates any Agentforce release decision. The measurement exists. The institutional response does not. That is the level-2-versus-level-4 gap, and it is the gap that the rest of this book is about.

The fix begins with vocabulary. A leadership team that uses the word “eval” precisely will start asking different questions. The first different question is the one you should ask in your next product review, and we will get to it in Chapter 8. For now, the work of this chapter is done. You have a definition. You have a taxonomy. You can listen to the next vendor pitch with new ears, and you will hear which level of seriousness is actually being claimed.

That alone is worth the next forty minutes.



Chapter 2 — How We Got Here: The Eval Economy in 18 Months

The inflection point

The eval economy was always coming, but it crystallized on a specific date. On June 13, 2025, Meta announced that it had taken a 49% stake in Scale AI for \$14.3 billion. Within hours, the procurement teams at OpenAI, Google, Anthropic, and xAI were on calls with their internal counsel asking the same question: can we still trust a vendor that is half-owned by a competing lab? Within weeks, the answer was visibly no. Google, which had reportedly planned to spend \$200 million on Scale AI services that year, materially reduced its engagement. OpenAI and Anthropic followed. The labs began calling the also-rans — Mercor, Surge, Turing, Handshake, Labelbox — and asking how fast they could ramp.

By October, Mercor had closed a \$350 million Series C at a \$10 billion valuation, quintupling from its \$2 billion price tag eight months earlier. Surge AI, which had been quietly profitable since founding, was reportedly raising at \$25 billion. LMArena, the academic spin-out from UC Berkeley that runs the most-watched LLM leaderboard in the world, raised \$150 million at \$1.7 billion in January 2026. Mechanize, six months old, was offering software engineers \$500,000 salaries to author reinforcement learning environments for Anthropic. Anthropic itself was reported to be considering more than \$1 billion in environment spending over the next twelve months.

This is what an industry reorganization looks like in real time. The category of “AI training data” — which most enterprise software leaders had filed under “commodity labor marketplace” — became, in eighteen months, a multi-billion-dollar layer of the AI stack with real moats, real margins, and real strategic significance. The companies winning that layer are quietly billion-dollar businesses that almost no consumer has heard of. They invoice in seven and eight figures per contract, sell to roughly twenty buyers globally, and have no consumer brand surface.

If you have not been paying attention to this category, you have missed the most important infrastructure shift in AI since the GPU bottleneck of 2023. This chapter is the catch-up.

Three forces converged

The first force was **data exhaustion**. By early 2024, the frontier labs had effectively trained on the entire publicly available text corpus of the internet. The next gains in model capability could not come from more text. They had to come from *better* signal — preference data, expert judgments, multi-step task traces, verifiable reward signals. That kind of data does not exist in scrapeable form. It has to be manufactured.

The second was **the failure of crowdsourcing at the new quality bar**. The previous generation of AI training data — image labels for autonomous vehicles, content moderation flags for social platforms — could be produced by crowdworkers earning \$3 per task. The new generation could not. Teaching a model to reason about a contract, debug a kernel, or diagnose a patient required *experts* doing the labeling: lawyers, kernel engineers, physicians. Mercor’s average contractor earns roughly \$85 per hour, and the platform’s daily contractor payouts run to about \$1.5 million. The economics of labor changed underneath the category, and the companies that adapted — Surge first, then Mercor — built around that shift while Scale, encumbered by its crowdworker DNA, struggled to catch up.

The third was **the saturation of static benchmarks**. By 2024, every credible LLM benchmark — MMLU, HumanEval, GSM8K — was either contaminated (the test set had leaked into training data) or saturated (every model scored above 90%). The labs needed live, contamination-resistant, human-preference-grounded evaluation infrastructure. LMArena, originally a UC Berkeley research project called Chatbot Arena, became the de facto neutral benchmark precisely because it was crowdsourced, dynamic, and academically governed. Once it was clear that the labs would pay for that evaluation infrastructure as much as they paid for training data, the third leg of the eval economy emerged.

The Meta-Scale deal of June 2025 was not the cause of any of this. It was the catalyst that converted a slow-building structural shift into an overnight market reorganization. The labs already wanted alternatives; the deal made the alternatives mandatory.

The three layers

The current eval economy resolves into three distinct layers, each with its own buyer profile, business model, and competitive dynamics.




Figure 2.1 — The eval economy in three layers. Frontier labs at the top buy from three distinct supplier categories: data factories, RL gyms, and eval arenas. Valuations and revenue figures as of Spring 2026.

The data factories sell expert-labeled training data, predominantly RLHF preference pairs. The buyer is a frontier lab post-training a model. The labor is architect-tier or domain-expert-tier humans — lawyers, doctors, software engineers — supplying preference judgments and rubric-graded scores. Mercor, Surge, and Scale are the dominant players. Valuations range from \$10 billion (Mercor) to \$25 billion (Surge, in current fundraising) to \$29 billion (Scale, post-Meta). Margins run high; Surge has been profitable since launch, which is the market’s tell that this is real infrastructure rather than dressed-up labor arbitrage.

The RL gyms sell sandboxed task environments — virtual machines where AI agents practice multi-step work and earn deterministic reward signals. The buyer is again a frontier lab, but specifically the team training the next generation of agentic models. The product is closer to a video game than a dataset: a coding sandbox with a verifiable correctness oracle, a browser environment with task completion checks, a customer-service simulator with conversation outcomes. Mechanize and Prime Intellect anchor this layer. Most companies are seed-stage

with under twenty employees, often serving just one or two lab customers under exclusive contracts. The category is fragmented and contested. The signal that it matters is Anthropic's reported \$1 billion+ environment commitment for the next year.

The eval arenas sell preference and judgment platforms, often with public leaderboards. The buyer is broader — frontier labs use them, but so do enterprises evaluating which model to deploy and developers choosing model APIs. LMArena is the dominant brand; Braintrust serves the enterprise EvalOps use case. The business model is consumption-based: pay-per-evaluation pricing scaled against model traffic. LMArena's annualized consumption run rate surpassed \$30 million in December 2025, less than four months after launching its enterprise product. Neutrality is the core moat — an arena that becomes captive to one lab loses its credibility, which is why LMArena's founders explicitly preserved academic governance after incorporating.

These three layers are not interchangeable. Each solves a different bottleneck — data exhaustion, agentic capability, contamination-resistant judgment — and the labs typically buy from all three in parallel. The reason consolidation has not happened yet, despite all the venture money chasing it, is that the layers genuinely require different operational competencies.

Why these companies exist at all

Here is the question worth lingering on: why did the labs not just do this work themselves?

OpenAI, Anthropic, and Google each have effectively unlimited capital and access to top engineering talent. They could, in principle, build their own annotator pools, design their own RL environments, host their own evaluation arenas. For some of this work, they do. But most of it, increasingly, they buy. The reason is structural and worth understanding.

Producing high-quality eval and training data at industrial scale is a different operational competency from training models. It requires recruiting and retaining tens of thousands of domain experts globally, managing payroll across jurisdictions, designing rubrics that calibrate consistently across raters, running quality-control workflows that catch low-effort submissions, and maintaining inter-rater reliability over multi-month engagements. That is closer to running a global temp staffing agency with PhD-level credentialing than it is to running an AI research lab. The labs are not built for it, do not want to be built for it, and accept paying a 30–50% margin to vendors who specialize in it.

This is the same logic that produced the SAP and Oracle implementation services boom of the 1990s, when Accenture, Sapient, and the IBM Global Services arm captured most of the economic value of the ERP wave despite producing none of the underlying software. The platforms could build the products; they could not, at the same time, scale the implementation

labor. The labor business became its own multi-billion-dollar industry adjacent to the platform business. The eval economy is an even cleaner version of the same pattern, because the “implementation” in this case is the calibrated human judgment that becomes a permanent input to model improvement.

NGP Capital, in October 2025, framed it bluntly: *EvalOps is no longer a cost center, it’s the differentiator*. What the labs are buying is not labor; it is calibrated judgment at industrial scale, packaged as preference data, environment episodes, or arena rankings. That is structurally a different good than annotation work, and it is why the valuations look more like infrastructure-tier multiples than like staffing-firm multiples.

What this means for enterprise platforms

There is a hard lesson in all of this for enterprise software companies, and it is the lesson the rest of this book is built around.

The frontier labs figured out, by late 2023, that eval and training data infrastructure was the rate-limiting input to model quality. They responded by building internal eval teams with release-gating authority, by signing nine-figure contracts with vendors who could supply expert judgment at scale, and by treating EvalOps as a first-class engineering discipline reporting to the CTO. The result is that the gap between the best frontier models and the second-best has narrowed dramatically over the past eighteen months — but the gap between the labs that took eval discipline seriously and the ones that did not has widened.

Enterprise software companies, with rare exceptions, did not get the memo. They have benchmarks. They have research papers. They have Trust Layers and guardrails and observability dashboards. What they do not have is the institutional muscle to translate measurement into release control, to translate user preference into training signal, to translate evaluation findings into roadmap reprioritization. The infrastructure of eval-driven development exists in the labs and is sold by the eval-economy vendors; it has not yet been installed in the enterprise platforms.

That is the gap Spring 2026 is about to expose. It is the gap Salesforce, specifically, is exposed to. And it is the gap Rosetta Phase 2 is designed to close.

The remaining six chapters explain how.



Chapter 3 — Anatomy of an Eval

Chapter 1 gave you the working definition: a structured measurement of model behavior against a defined rubric, with verifiable scoring, used to drive product decisions. Chapter 2 explained why the eval economy emerged. This chapter is the technical primer that lets you read a vendor pitch, evaluate a research paper, or sign off on an internal eval design without faking the technical literacy.

We have aimed it at the smart non-engineer. The technical concepts are introduced through a consistent worked example and grounded in the academic literature where the field has settled. If you are an engineer who already builds evals for a living, skim this chapter; the next one is more useful to you.

The seven building blocks

Every eval, from the simplest LLM-as-judge call to the most elaborate multi-step RL gym episode, is composed of the same seven elements. Naming them precisely is the foundation for everything that follows.

1. The input prompt. The question, instruction, or scenario that the system under evaluation receives. Specified in writing, with no ambiguity about formatting or framing. *“What is our refund policy?”* is an inadequate prompt; *“User asks: ‘What is our refund policy for products purchased on sale during the December 2025 promotion?’ Customer ID 4471, account in good standing, prior return history attached”* is an adequate one.

2. The context window. What other information is provided alongside the prompt — system instructions, retrieved documents, conversation history, tool definitions. The context is part of the test. An eval that varies the context across runs is testing different things and producing non-comparable results.

3. Candidate response(s). What the system produces. For preference evals, you collect two or more candidate responses — often from different model versions, retrieval configurations, or prompts — so the grader can compare. For capability evals, you typically collect a single response and judge it against a target.

4. The rubric. The written specification of what counts as a good response. The rubric is the hardest artifact in the entire eval program. We will return to this in detail below.

5. The verifier. The mechanism that converts a response into a score. Could be deterministic (does the JSON parse correctly?), an LLM acting as judge (does this response satisfy the rubric

criteria?), or a human expert (preference judgment between candidates A and B). The next section walks through the trade-offs.

6. The score. The numeric or categorical output of the verifier, plus structured metadata: which rubric criterion drove the score, how confident the verifier was, where on the scale the response landed.

7. The justification. A written explanation of why the score is what it is. Often skipped by inexperienced eval programs and always regretted later. Without justifications, you cannot audit eval quality, you cannot retrain raters who are drifting, and you cannot defend results when stakeholders push back.

The Press (2026) academic formulation captures this as a four-tuple: a benchmark consists of *Request*, *Environment*, *Stopping Criteria*, and *Scorer*. Our seven elements decompose those four further by separating context from input, response from candidates, and score from justification. Same idea, more granularity for practitioners.

The three verifier types

The verifier is where most of the cost and most of the methodological argument live. There are three types, and the right choice depends on what you are measuring.

Deterministic verifiers are programs. They check whether a response satisfies a precisely specified condition: does the SQL parse, does the function pass its unit tests, does the answer contain the required citation. Cheap to run (microseconds, fractions of a cent), perfectly reproducible, and correct by construction within their narrow scope. The catch: they only work on tasks where correctness can be specified in code. For most real-world enterprise scenarios, that is a small fraction.

LLM-as-judge verifiers are language models prompted with a rubric and asked to score a response. Cost per evaluation is roughly \$0.01–\$1.00 depending on model and rubric complexity. The advantages are speed and scale; you can run a million of them overnight. The disadvantages are calibration (LLM judges drift, develop subtle biases, and disagree with human experts on edge cases) and contamination risk (the judge model may have been trained on data that influences its preferences). Used carefully — typically as a first-pass screen with human spot-check on a calibration set — they are the workhorse of modern eval programs.

Human expert verifiers are domain experts judging responses against the rubric. Cost is hourly: \$50 for crowdworkers on simple tasks, \$200–\$500 for SME-tier judgments (lawyers, doctors, software architects). Slow, expensive, and the gold standard for any subjective task where human preference is the ground truth. The frontier labs spend the majority of their eval budget here for exactly this reason.

In practice, serious eval programs use **hybrid verifier strategies** that combine all three: deterministic checks catch obvious errors, LLM judges screen the majority of responses, and human experts grade the hardest 5–20% and calibrate the LLM judges against their own judgment. The Phase 2 Rosetta engagement is structured exactly this way.

The rubric is the hard part

Most eval programs that fail, fail at the rubric.

Writing test cases is engineering work; writing rubrics is judgment work. A good rubric specifies, in writing, what counts as a correct or preferred response across all the dimensions that matter for the use case. It must be precise enough that two trained graders applying it to the same response will reach the same score most of the time. It must be complete enough that the criteria it omits are not the ones that actually matter to users. And it must be stable enough that the rubric you use in week 12 is comparable to the rubric you used in week 1, or the program's results lose continuity.

This is genuinely difficult. A novice rubric specifies “the response should be helpful” and discovers, three weeks later, that its graders are interpreting “helpful” in mutually incompatible ways. A mature rubric breaks “helpful” into four to six measurable sub-criteria, each with explicit examples of what passes and fails, and includes a calibration set of twenty to fifty reference responses with documented scores so new graders can be trained against the same standard.

The rubric is the unit of disagreement in an eval program. When two graders disagree, the rubric is what you go back to. When the eval results disagree with what the product team believed, the rubric is what you defend or revise. When a competitor publishes a benchmark you do not believe, the rubric is what you critique. Any time the program is invested in writing rubrics, in our experience, returns roughly 5x in downstream stability — fewer re-grading cycles, fewer methodological disputes, faster onboarding of new graders.

Cohen's kappa and what 0.7 means

The standard metric for measuring whether your graders agree with each other is **Cohen's kappa**, a statistic that ranges from 0 (graders are no more aligned than random chance) to 1 (perfect agreement). Kappa is preferred over raw agreement because it adjusts for the agreement you would expect by accident on imbalanced rubrics.

The conventional thresholds:

- **$\kappa < 0.40$** — Poor agreement. The rubric is broken, the graders are undertrained, or both. Eval results should be discarded.

- $0.40 \leq \kappa < 0.60$ — Moderate agreement. Useful for exploratory work; not reliable enough for training data or release decisions.
- $0.60 \leq \kappa < 0.80$ — Substantial agreement. The threshold for production-quality eval programs. Most published RLHF work targets the upper half of this range.
- $\kappa \geq 0.80$ — Near-perfect agreement. Possible on highly objective tasks; rare on the kinds of subjective judgments that actually require human raters.

The widely used 0.7 floor for RLHF preference data sits in the middle of the “substantial” band. A program reporting kappa above 0.7 is producing eval data that can drive training and release decisions. A program reporting kappa below that — or a program that does not report kappa at all — is producing data that should not be used for either purpose. When you are reading a vendor’s eval claims, the absence of an inter-rater reliability metric is itself a signal.




Figure 3.1 — Calibration drift over twelve months. Three program types compared on a fixed reference set. Illustrative; based on observed patterns in long-running eval programs.

The three drift modes

Long-running eval programs degrade in three predictable ways. Each requires a different countermeasure.

Rubric drift is graders interpreting the same rubric differently as time passes, often because they encounter edge cases the rubric did not anticipate and resolve them inconsistently. The countermeasure is quarterly recalibration: a fixed reference set of responses re-graded by all active raters, with kappa recomputed and the rubric revised when agreement falls below threshold.

Model drift is the system under evaluation changing in ways that invalidate prior eval results. A model retrained on new data may answer the same prompt differently; a retrieval system reconfigured may return different context. The countermeasure is versioning every component of the eval and treating cross-version comparisons as a separate analytical step rather than a default assumption.

Domain drift is the world changing underneath the eval. Customer language evolves, product features ship, regulations update. An eval that was valid in Q1 may be measuring the wrong thing by Q4. The countermeasure is periodic ground-truth refresh — a structured review of whether the rubric and reference scenarios still reflect what users actually care about. Most enterprise eval programs underestimate this and discover the gap only when results stop correlating with user satisfaction.

The OpenReward Standard (General Reasoning, 2026) explicitly addresses these drift modes by serving evals as managed API endpoints with versioned environments and curriculum management, allowing 330+ RL environments backed by 4.5 million tasks to remain coherent across model generations. Enterprise programs without that infrastructure tend to rebuild evals from scratch every six to nine months, at considerable cost.

A worked example: identity resolution

Here is one Rosetta eval, walked through end-to-end, to put the seven elements together.

The scenario is an identity resolution question in Data 360. A Salesforce administrator is consolidating customer records after an acquisition. Two records exist: one in the legacy CRM as “Robert Chen,” account ID 4471, with a Gmail address; one in the acquired company’s system as “Bob Chen,” account ID A-9982, with a corporate email at the acquired entity. The administrator asks: *“Are these the same person? If yes, which record should be the survivor?”*

- **Input:** the prompt above.
- **Context:** both account records, transaction history for each over twenty-four months, and the acquired company’s identity resolution policy document.
- **Candidate A:** Rosetta’s current production response — confidently asserts they are the same person, recommends keeping the older legacy record as survivor, does not flag the ambiguity in the email-domain transition.
- **Candidate B:** Rosetta with the proposed retrieval improvement — flags this as a probable match (estimated 87% confidence) based on three converging signals (name, transaction overlap, IP geolocation), recommends a human review step before merge, and surfaces the email-domain ambiguity as a specific risk.

- **Rubric:** five criteria — *factual correctness, appropriate confidence calibration, governance compliance* (does it suggest a process the customer’s data stewards would endorse), *use of customer vocabulary*, and *surfacing of risks the user did not ask about*.
- **Verifier:** an LLM judge runs first-pass screening against the rubric; a human expert (D360-experienced solution architect) grades 100 of these 1,000 pairs against the rubric and provides written justifications, with a second architect re-grading 10% of the human-judged pairs to compute inter-rater reliability.
- **Score:** Candidate B wins on four of five criteria (loses on conciseness, which was not part of this rubric); preferred by the human grader 8.4 times out of 10; LLM-judge agreement with human grader on the 100-pair calibration set yields $\kappa = 0.74$.
- **Justification (sample, written by the human grader):** *“Candidate B correctly identifies match probability as below the auto-merge threshold, surfaces the email-domain change as a governance risk the user should consider before deciding, and recommends a process step (human review) consistent with the customer’s documented identity resolution policy. Candidate A’s confident assertion is technically defensible on the data but procedurally inappropriate; in production, it would lead to merge errors that customer data stewards would later have to repair.”*
- **Decision:** Candidate B’s retrieval improvement ships, gated on a regression test ensuring future model versions maintain $\geq 85\%$ preference rate against Candidate A’s behavior on this scenario family.

That is one eval. It contains all seven building blocks, uses two of the three verifier types, reports an inter-rater reliability metric above the 0.7 threshold, and produces a release decision. It is the unit of work the rest of this book is about scaling.

Closing

You now have the technical vocabulary to read any eval result with literacy: the seven elements, the three verifier types, the rubric as the hard artifact, Cohen’s kappa as the reliability metric, the three drift modes as the failure modes. The next chapter takes these primitives and organizes them into a working taxonomy of the five eval types you will encounter in enterprise AI conversations — and the one type, almost always missing, that is the cause of most enterprise AI quality failures in the field.



Chapter 4 — The Five Eval Types You Need to Know

A taxonomy is useful when it lets you put a name to a thing you previously could not distinguish. The five-type taxonomy in this chapter does that for the field of AI evals. It is not the only taxonomy in circulation; academic researchers slice the territory differently, and so do the major eval-economy vendors. But these five categories, in our experience, are the ones a Salesforce executive or architect needs in working memory to navigate the conversations of Spring 2026 without being talked past.

We will define each type, give a representative public benchmark or vendor offering, and place it on a two-axis map. The chapter closes by naming where Rosetta sits and where the highest-leverage gap is in current enterprise AI practice.

The five types

1. Capability evals. These ask: *can the model do X at all?* The classic public benchmarks live here — MMLU for general knowledge, HumanEval for code, GSM8K for grade-school math. In the enterprise CRM context, Salesforce AI Research’s CRMArena and SCUBA are capability evals: structured task sets that score whether the model successfully completes representative work. Capability evals are the first thing built in any new AI domain because they are the easiest to design and the easiest to communicate to executives. They are also the easiest to game, the easiest to contaminate, and the most likely to saturate. A model scoring 95% on a capability benchmark you designed two years ago is almost certainly telling you something about contamination, not about capability.

2. Reliability evals. These ask: *does the model do X consistently across edge cases?* Where capability evals measure peak performance on representative scenarios, reliability evals measure failure rates across the long tail. They are explicitly designed to find the cases where the model breaks: ambiguous inputs, adversarial phrasings, edge cases the training distribution underweighted, scenarios where two parts of the rubric come into tension. Anthropic’s red-team evals fall in this category. Reliability evals are systematically underbuilt in enterprise platforms because they are politically uncomfortable — they exist to find problems, and finding problems creates work. The mature posture is to fund them generously precisely because no one else will.

3. Preference / alignment evals. These ask: *between two reasonable responses, which does the human expert prefer, and why?* This is the RLHF data category. Preference evals are the

only honest answer for tasks where there is no single right response — most real-world enterprise scenarios. The output is graded preference pairs with rubric-justified judgments, used both to evaluate current model behavior and to train future model behavior. This is what Phase 2 of the Rosetta program is producing. It is also what the data factory companies — Mercor, Surge — sell, and what the frontier labs spend the majority of their eval budget on.

4. RL gym evals. These ask: *can the agent execute a multi-step task successfully against a deterministic verifier?* RL gyms are sandboxed environments where an agent attempts a task — write a function, complete a customer-service interaction, navigate a software UI — and the environment provides verifiable success or failure feedback. Mechanize and Prime Intellect sell into this category. Salesforce’s SCUBA, despite its name as a benchmark, is structurally an RL gym: a simulated CRM environment where computer-use agents attempt realistic tasks and a verifier checks completion. RL gym evals are the highest-fidelity measurement available for agentic capability, and they are also the most expensive to build, because each environment is essentially a custom piece of software.

5. Production evals. These ask: *is the model still working as expected, in the wild, on live traffic?* Production evals run continuously against the actual user requests hitting the deployed system, not against a frozen test set. They detect drift (model behavior changing over time), regression (a new release degrading on cases the old release handled), and emergent failure modes (categories of failure that did not exist in the test set because no one had thought of them yet). This is the category almost every enterprise platform is missing in Spring 2026. It is also the category where the highest customer-facing leverage exists, because production evals are the only kind that catch the failures customers actually experience.

A 2x2 to keep these straight

The five types organize cleanly along two axes. The first axis is *deployment lifecycle stage*: pre-release versus post-release. The second axis is *what the eval is doing*: measurement (telling you how the system performs) versus control (gating a decision based on the result).




Figure 4.1 — The five eval types on a 2x2 of measurement vs control by pre- vs post-release. Two types (preference and production evals) span the measurement-control boundary because they both score and gate.

The pattern visible in this matrix is the diagnostic for an enterprise platform's eval maturity. A platform that has filled the pre-release × measurement quadrant has done the work that gets published in research papers. A platform that has filled the post-release × control quadrant has done the work that protects customers from quality failures in production. The frontier labs have filled both. Most enterprise platforms have filled the first and almost none of the second. The remainder of this book is largely about the consequence of that asymmetry.

Where Rosetta sits

The Rosetta program, in its Phase 2 design, is concentrated in the preference eval category — the third type — with deliberate extensions into production eval as the program matures. This is a considered choice. We will explain the reasoning briefly here and return to it in Chapter 7.

Capability and reliability evals are well-supplied by the existing Salesforce AI Research work. CRMArena, ContextualJudgeBench, and SCUBA are credible benchmarks; the technical work to build more in those categories is not where the highest marginal value sits. RL gym evals are extremely expensive to build well and are the contested territory of the Mechanize / Prime Intellect tier; for a knowledge-graph-centered program like Rosetta, building a gym from scratch is not the right early bet.

Preference evals, by contrast, are exactly where Rosetta has a structural advantage. The canonical inventory underlying the Rosetta knowledge graph is a rare and defensible source of ground truth. The program's relationship with practitioner architects — Ian Gotts, Mehmet Orun, the broader Salesforce architect community — is the access path to expert graders that Mercor would have to build from scratch. The combination produces preference data that is more grounded in actual customer scenarios than what generic SME pools can supply, and that compounds in value as the canonical inventory grows.

Production evals are the obvious extension. Once the rubric IP, the verifier infrastructure, and the calibrated graders are in place from Phase 2, running them continuously against live D360 traffic is a natural Phase 3 deliverable. That progression — from preference evals as the foundational asset to production evals as the customer-facing instrumentation — is the full arc of what an EvalOps program for Salesforce should look like.

The missing type

The chapter ends with a diagnostic claim that sets up the rest of the book.

The category most missing from enterprise AI in Spring 2026 is the fifth: production evals. The first four types are visible in research papers, vendor pitches, and benchmark leaderboards. The fifth is invisible because it does not produce publishable artifacts. It produces operational signal — drift detection, regression alerts, quality dashboards — that lives inside the engineering organization and never surfaces externally unless something goes wrong.

The result is that enterprise platforms can plausibly claim eval maturity on the basis of categories 1 through 4 while being almost entirely undefended on category 5. This is not hypothetical. The Agentforce hallucination variance documented in field reports — somewhere between 3% and 27% depending on configuration — is a production eval failure mode. A platform with mature production evals would know which customers are sitting at 27%, would have detected the drift before customer complaints surfaced it, and would have the instrumentation to gate further deployment until the configuration issue was corrected. Most platforms in Spring 2026 do not have that instrumentation.

This is the gap Chapter 6 documents at Salesforce specifically. It is the gap Chapter 7 explains how Rosetta closes. For now, the work of this chapter is done: you can name the five types, place any eval claim you encounter on the 2x2, and recognize that the type most likely to be missing is the one that protects customers. That recognition alone is enough to ask the right question in your next product review.

The next chapter is the one you read if you want to know what taking these five categories seriously actually looks like inside the organizations that have been doing this for two years longer than the rest of us.



Chapter 5 — How the Best AI Companies Actually Use Evals

This chapter describes practice, not theory. What we report here is how Anthropic, OpenAI, Google DeepMind, and the other frontier labs actually operationalize eval-driven development in Spring 2026, drawn from public commitments, published research, hiring patterns, vendor relationships, and conversations with practitioners inside and adjacent to those organizations. We have stripped the marketing veneer and tried to describe the operating reality.

The reason this chapter exists is that “we do evals” means something very different at Anthropic than it does at the median enterprise software company. Until you can see the difference clearly, the gap documented in Chapter 6 reads as criticism rather than as diagnosis. Our intent is the latter.

The eval-gated release principle

The single most important practice the frontier labs have institutionalized — and the one most absent from enterprise platforms — is **eval-gated release**. Models do not ship until they pass internal eval thresholds. Eval thresholds are set in advance. Eval results have authority to block release. The team running the evals is organizationally separate from the team building the model, and they can say no.

Anthropic’s Responsible Scaling Policy, public since 2023 and repeatedly updated, is the cleanest published example. The policy commits the company to specific evaluation protocols at defined capability thresholds, with explicit deployment and security commitments tied to evaluation outcomes. The policy is enforceable internally — Anthropic has on at least one occasion delayed a release publicly because of evaluation findings — and it is the primary instrument that translates eval results into product decisions.

OpenAI’s Preparedness Framework operates similarly, with capability thresholds across four risk categories and pre-defined safeguards triggered by evaluation outcomes. Google DeepMind’s Frontier Safety Framework, published in 2024, takes the same shape: defined

capability levels, evaluation protocols at each level, deployment commitments contingent on results. The vocabulary differs across labs; the structural pattern does not.

Compare this to the typical enterprise software release process. A model or feature is ready when the engineering team says it is, with informal sign-off from product management and a final QA pass against a written specification. There is no organizational entity with the authority to block release based on evaluation findings, because the evaluation function does not exist as an independent organizational unit. This is the asymmetry that defines the gap.

The lesson is not that enterprise platforms need to publish a Responsible Scaling Policy. It is that enterprise platforms need an internal function with the authority to say no — and the eval data infrastructure to back up the no with evidence. Without both, “we do evals” is performative.

The constitutional approach to rubric design

Anthropic’s Constitutional AI methodology, published in 2022 and refined since, is interesting beyond its training implications because of what it reveals about how the company approaches rubric design. The “constitution” is, effectively, a written rubric: a specification in natural language of what counts as a desirable response across a set of criteria. The model is trained to evaluate its own responses against the rubric and revise them when they fail.

The deeper lesson for enterprise practitioners is the philosophical one: Anthropic treats the rubric as the durable artifact and the model as the perishable one. Rubrics are versioned, debated, refined over time. Models come and go; the constitution persists across model generations. This is the opposite of how most enterprise platforms relate to their evaluation infrastructure, which treats the model as the durable asset and the test cases as one-off project artifacts that are rebuilt every release cycle.

The practical implication is that frontier labs invest heavily — disproportionately, by enterprise-software-engineering standards — in writing, calibrating, and maintaining rubrics. A rubric that has been refined over twenty thousand grading decisions is more valuable than a model checkpoint, because the rubric encodes accumulated judgment about what good looks like, and that judgment is the rate-limiting input to every future model improvement.

Enterprise platforms that want to install eval discipline should start by asking: who, in this organization, is the rubric author of record, and where is their rubric stored? In most cases, neither question has an answer, because the role does not exist. That is itself a diagnostic.

The preference data flywheel

OpenAI's relationship with preference data illustrates the second pattern: the flywheel. Models in production generate responses. Users rate or revise those responses. The rated data, after careful filtering and quality control, becomes training signal for the next generation of models. The cycle compounds — each iteration produces a model that generates higher-quality responses, which means user feedback is more discriminating, which means the next iteration's training signal is more valuable.

This pattern is harder to import into enterprise platforms than it appears. The difficulty is not technical; it is contractual and cultural. Customer enterprises do not typically grant their software vendors the right to use customer interactions as training data, with good reason. Without a structured program to acquire that right — usually involving anonymization, opt-in, and occasionally direct compensation — the flywheel does not turn.

Mercor's emergence as a \$10 billion company is in part a response to this constraint. The labs cannot acquire preference data from production traffic at the scale they need; they buy it from a vendor that pays experts to generate it synthetically against scenario specifications the labs supply. Surge AI's \$1+ billion in 2024 revenue tells the same story. The labs are buying what they cannot manufacture in-house.

Enterprise platforms have a structural advantage here that they have largely failed to use. They have direct relationships with the customers whose data and judgments would be most valuable as eval inputs. A Salesforce that systematically engaged its customer architects as preference graders — with appropriate compensation and IP arrangements — could build a preference data corpus that no frontier lab could replicate, because the labs do not have those relationships. Phase 2 of the Rosetta program is, in part, a small-scale test of exactly this thesis.

The Mercor lesson: experts, not crowdworkers

The single most important methodological shift in the eval economy over the last eighteen months is the move from crowdsourced labeling to expert labeling. Scale AI's original business was built on hundreds of thousands of crowdworkers labeling autonomous-vehicle bounding boxes for three to five dollars per task. Mercor and Surge built deliberately the opposite: smaller, higher-paid teams of domain experts. Mercor's 30,000 contractors average \$85 per hour; the platform pays out roughly \$1.5 million per day. Surge has been profitable since launch, which on labor-arbitrage economics would be impossible.

The reason the shift happened is that the new generation of evals requires judgments crowdworkers cannot make. Asking a crowdworker to evaluate whether a legal opinion correctly cites case law, or whether a kernel patch is idiomatic, or whether a sales response addresses

the customer's actual underlying objection, produces noise rather than signal. The signal requires someone who could write the response themselves at the same quality level as the model under evaluation.

For enterprise eval programs, the implication is direct: do not start with a crowdsourced grading pool. Start with a small number of vetted domain experts and accept that the cost per eval will be five to ten times higher than a generic SME pool. The Phase 2 Rosetta engagement is structured this way — architect-tier graders at \$300 per hour loaded — because that is the only way to produce eval data that holds up under technical scrutiny. Trying to do this on the cheap produces eval data that fails inter-rater reliability checks and is then unusable for any serious purpose.

The Surge lesson: rubric IP, not labor margin

Surge's profitability is the tell. A pure labor-marketplace business at SME tier should run on thin margins; the contractor takes most of the revenue and the platform skims a coordination fee. Surge does not operate that way. The company invests heavily in proprietary rubric design, custom verifier infrastructure, and quality-control workflows that compound across engagements. What the labs are paying for, when they pay Surge, is not access to expert labelers — they could find those independently — but access to Surge's accumulated methodology for turning expert judgment into reliable, repeatable, training-grade data.

This is the durable moat in the eval economy. Labor is reproducible; rubric IP is not. A platform that has refined a rubric across thirty engagements has a measurement instrument that a new entrant cannot match without paying the same five-year tuition. The same logic applies to verifier libraries, calibration sets, and inter-rater training protocols. These are the artifacts that make eval businesses durable, and they are exactly what enterprise platforms tend not to build because the procurement function does not know to ask for them.

For Salesforce specifically, the implication is that buying eval services from a vendor produces eval data; commissioning the development of rubric IP and verifier infrastructure as a deliverable produces a durable internal asset. The two engagements look superficially similar and have radically different long-term value. Phase 2 of the Rosetta program is structured around the second model.

The LMArena lesson: neutrality is the moat

LMArena's \$1.7 billion Series A in January 2026 is impossible to explain based on revenue alone. The company's annualized consumption run rate at the time of the raise was

approximately \$30 million. The valuation makes sense only if you understand what LMArena is selling: trust.

Every AI model vendor publishes self-reported benchmark numbers. Every enterprise customer evaluating those vendors has learned to discount self-reported numbers heavily. What does not exist, with rare exceptions, is a neutral third party where the same models are evaluated on the same tasks with the same methodology, producing comparable results that none of the vendors can manipulate. LMArena is that party. The leaderboard is credible because the operating organization has explicit governance commitments preventing any single lab from controlling the methodology or the results. Neutrality is the entire product.

The lesson for enterprise platforms is uncomfortable. A platform vendor can build internal evals; what it cannot build, by definition, is neutral evals about itself. Customers and analysts will treat vendor-published eval claims with appropriate skepticism, because the vendor is not a neutral party. The only way to produce credible third-party-grade evidence about a platform's quality is to subject it to evaluation by an organization with explicit independence from the platform.

Salesforce has no equivalent. There is no LMArena for CRM. The first organization to build one — under credible neutrality covenants — captures the same trust premium that LMArena now commands in the LLM market. This is part of what Phase 3 of the Rosetta program contemplates, and we will return to it in Chapter 7.

The org chart tells the story

If you want a single diagnostic for whether an organization has installed eval discipline, look at the org chart.

At the frontier labs, evaluation teams are independent organizational units with named leaders, budgeted headcount, dedicated infrastructure, and reporting lines that go to the CTO or directly to the executive team. Anthropic's evaluation team operates with explicit independence from the model training organization. OpenAI's safety and evaluations function reports outside the product line. Google DeepMind has a dedicated Frontier Safety team. The structural pattern is consistent: evaluation has its own org, its own budget, and its own authority.

At the typical enterprise software company in Spring 2026, evaluation responsibilities are distributed: some sits in QA, some in data science, some in the product team running the AI feature, none of it in an organization with the authority to block release. The hiring pattern follows the org structure: instead of evaluation engineers and rubric stewards, the open requisitions are for "Forward Deployed Engineers" who troubleshoot failures after deployment. The hiring posture is firefighting, not prevention.

This is the most actionable diagnostic in this entire book. If you want to know whether your organization has installed eval discipline, do not ask whether your team is doing evals. Ask: where does the evaluation function report, and what authority does it have to delay a release? If the honest answer is “nowhere” and “none,” then “we do evals” is aspirational, regardless of how many benchmark scores appear in the next quarterly review.




Figure 5.1 — Two org charts, one diagnostic. The frontier-lab pattern installs an independent Evaluation function with release-gating authority. The typical enterprise-SaaS pattern distributes eval responsibilities across peers and engages Forward Deployed Engineers reactively after deployment.

What to take from this chapter

The frontier labs are not magical. They are not staffed exclusively by geniuses, they do not have access to secret evaluation methodologies, and the practices described in this chapter are within reach of any well-run engineering organization. What they have done is install the institutional infrastructure that converts measurement into decision authority and accumulated rubric work into durable competitive advantage. The investment is real — eval teams of dozens to hundreds of people, vendor commitments in the nine-figure range — but it is not exotic.

The gap between the labs and the enterprise platforms in Spring 2026 is therefore not a gap of capability. It is a gap of priority and organizational design. Enterprise platforms that decide to close the gap can close it within a release cycle or two. The decision to close it has to be made by someone with the authority to allocate budget, name a function, and hand the function the

keys to the release process. That decision has not been made at most enterprise platforms. It has not yet been made at Salesforce.

The next chapter is about what that gap looks like at Salesforce specifically, in Spring 2026, with the data we have and the honesty the situation deserves.



Chapter 6 — Where Salesforce Is, Honestly

This chapter is the one we wrote the rest of the book to make readable.

The previous five chapters built the vocabulary, the taxonomy, the technical primitives, and the frontier-lab benchmark. With those in hand, the question of where Salesforce stands in Spring 2026 becomes answerable in concrete terms rather than as opinion. The answer is not flattering, and we are not going to soften it. Salesforce has the research talent, the customer footprint, and the institutional reach to lead enterprise AI evaluation. It has not built the cultural and organizational infrastructure to do so. The gap is real, it is widening, and it is closeable — but only on a deliberate schedule, starting now.

Three things before we begin. First, the diagnosis is not the people. The Salesforce AI Research team is genuinely capable, and several individuals within it would be welcome contributors at any frontier lab. The gap is structural, not personal. Second, we are reporting from public evidence: published research, the CRM Benchmark dashboard, job postings, customer field reports, and the internal program documents we have access to as part of the Foundation Sprint. We do not rely on confidential information to make any claim in this chapter. Third, we say this as advisors who have been inside the building. Truth-to-power is the Keenan Vision posture. The reason we can be useful to Marc and Steve is that we say what is true even when it is awkward. This chapter is the truth, as we see it, in Spring 2026.

What Salesforce has done well

Credit first, because the credit is real.

Salesforce AI Research has shipped serious public eval work. **CRMArena**, published in late 2024 and refined since, is a benchmarking framework that simulates realistic CRM scenarios and scores agent performance against them. **ContextualJudgeBench** evaluates response pairs across more than two thousand challenging cases, scoring accuracy, conciseness, faithfulness, and appropriate refusal. **SCUBA**, published on arXiv in October 2025, is a computer-use benchmark that tests agents on enterprise software tasks; the public results show

that even closed-source frontier models complete only 39% of zero-shot tasks and 50% with demonstration augmentation. The **CRM Benchmark dashboard**, published as a Tableau workbook, evaluates agents on accuracy, cost, speed, trust and safety, and sustainability across enterprise use cases.

This is the work of an organization that takes capability evaluation seriously at the research level. The team that produces it is, to put it simply, doing good work. If the rest of the company treated this output the way frontier labs treat their evaluation outputs, this would be a different chapter.

It does not, and so we proceed.

The first gap: measurement without gating

The most important diagnostic in this chapter takes one sentence to state.

Salesforce measures. Salesforce does not gate.

The SCUBA benchmark, by Salesforce's own published results, shows that the best available agentic models complete fewer than half of realistic CRM computer-use tasks. There is no public evidence — and no evidence we have been able to surface in our work inside the program — that this number gates any Agentforce release decision. There is no documented internal pass threshold on SCUBA. There is no organizational unit with the authority to delay an Agentforce release because of SCUBA scores. There is no public commitment, on the model of Anthropic's Responsible Scaling Policy or OpenAI's Preparedness Framework, that ties Salesforce's deployment decisions to specific evaluation outcomes.

The measurement exists. The institutional response does not.

This is the level-2-versus-level-4 gap from Chapter 1, and it is the gap that defines the rest of the diagnosis. Capability benchmarks that do not gate releases are research artifacts, not engineering controls. They produce papers and conference talks. They do not produce trustworthy products. The frontier labs figured this out by 2023 and built the institutional infrastructure to translate measurement into release authority. Salesforce, in Spring 2026, has not.

The test we proposed in Chapter 5 applies here directly. Convene a meeting with the Agentforce product team and ask: *what is our internal pass threshold on SCUBA, and who has authority to block a release on it?* Notice whether anyone can answer. If no one can, the gap is real, and no amount of additional benchmark work will close it without addressing the institutional structure.

The second gap: internal-facing, not customer-facing

The CRM Benchmark dashboard exists. Customers cannot use it.

This is a specific claim, and it is worth being precise about what we mean. There is no customer-facing tool today by which a Salesforce administrator can load their organization's specific Agentforce configuration, run a battery of evaluations against it, and see how their deployment scores against published benchmarks. There is no tool by which they can detect drift in their agent's behavior over time. There is no tool by which they can compare their deployment against industry baselines for hallucination rate, task completion, or response quality. The eval infrastructure is built; it is not productized.

The consequence shows up in the field. Hallucination rates in production Agentforce deployments span an order of magnitude — somewhere between 3% and 27% depending on configuration discipline — and customers cannot tell which side of the variance they are on except by way of customer complaints. By the time complaints surface, the customer's trust in the platform has already been damaged. Reviews on G2 and other platforms capture the pattern: enterprise users discovering, through production failures rather than instrumentation, that their Agentforce deployment is hallucinating product features and pricing details. *"The hallucination is really bad, since we don't train, and it works on a general model, sometimes it just gives information that isn't ours,"* one reviewer wrote in late 2025. That review should not have been necessary. Customer-facing evaluation tooling, deployed properly, would have caught the issue before the customer did.

This is the productization gap. The research team has built the measurement infrastructure. The product organization has not exposed it to the customers who need it most. The reason is not technical — the dashboard, the benchmarks, and the evaluation methodology already exist internally. The reason is that nobody owns the productization mandate. There is no Customer-Facing Eval Tooling product line, no roadmap commitment to deliver one, and no executive sponsor making the case that this should exist. Without that ownership, internal eval infrastructure stays internal, and customers continue to discover quality issues by way of failure.

The third gap: vendor-controlled, not ecosystem-shared

Every credible eval result Salesforce has published is on Salesforce-controlled infrastructure, scored by Salesforce-defined methodology, against Salesforce-curated test sets. We are not suggesting bad faith; the work is methodologically sound. We are pointing out the structural feature: there is no LMArena equivalent for CRM. There is no neutral third-party platform where Agentforce, Microsoft Copilot Studio, ServiceNow's AI agents, and the hyperscaler-native MCP

implementations are evaluated head-to-head on shared tasks under shared methodology, with results that none of the vendors can manipulate.

This is the neutrality gap, and it is the one that exposes Salesforce to a specific competitive risk. The first vendor in the CRM agent category to publish credible third-party-evaluated results — particularly results that compare favorably to Salesforce on tasks customers care about — captures a procurement narrative that Salesforce has no comparable public data to defend. The procurement conversation tilts overnight: not because Salesforce’s product became worse, but because a competitor produced evidence Salesforce does not have.

This is exactly what happened to Scale AI when Meta took its 49% stake. Scale’s product was unchanged the day after the announcement. Its competitive position was not. The trust premium Scale had built over years collapsed in weeks because the structural conditions under which that trust was credible had changed. Salesforce’s eval position is structurally similar: trustworthy as long as no neutral comparator exists, exposed the moment one does.

The defense is not to publish more vendor-controlled benchmarks. It is to participate in — or, more strategically, to *create* — neutral third-party evaluation infrastructure for the CRM agent category. The first organization to build that infrastructure under credible neutrality covenants captures the same trust premium LMArena now commands in the LLM market. Phase 3 of the Rosetta program contemplates exactly this, and we will return to it in Chapter 7. For now, the diagnostic point is that the neutrality gap is the most strategically dangerous of the three because closing it requires institutional commitments — independence covenants, methodology governance, ecosystem participation — that are easier to make early than late.

The Trust Layer is the wrong instrument

A note on the Trust Layer, because it comes up in every conversation about Salesforce AI quality and it is being used to do work it cannot do.

The Trust Layer, as Salesforce documents it, includes guardrails such as toxicity detection, PII redaction, prompt injection detection, and grounding mechanisms that tie LLM outputs to company-specific data. These are real and useful safeguards. They are also entirely downstream. The Trust Layer measures whether the agent did something *bad* — produced toxic output, leaked PII, was manipulated by a prompt injection. It does not measure whether the agent did the business task *correctly*.

These are different problems. A model can pass every Trust Layer check and still hallucinate the customer’s pricing structure. A model can fail Trust Layer checks and still answer the customer’s actual question competently. The Trust Layer is necessary; it is not sufficient. Conflating downstream safety with upstream quality measurement is a category error, and the

conflation appears repeatedly in Salesforce marketing materials in ways that obscure the gap this chapter is documenting.

This matters because customers and prospects, reading those materials, come away believing Salesforce has installed eval discipline when in fact Salesforce has installed safety guardrails. The two are not interchangeable, and the gap between them is exactly where the customer-facing failures are accumulating.

The hiring posture as evidence

Where an organization is investing tells you what it actually believes. The Salesforce hiring posture in Spring 2026 is unmistakable.

Open Salesforce job requisitions in the AI organization disproportionately ask for **Forward Deployed Engineers** charged with troubleshooting agent behaviors after deployment — diagnosing hallucinations, grounding gaps, orchestration failures, and integration errors in customer environments. The role is real, the work is necessary, and the people doing it are competent. But the posture the role represents is *firefighting*: post-deployment remediation of failures that should have been caught before deployment.

What is conspicuously absent from the hiring pattern is the alternative posture: **eval engineers, rubric stewards, calibration owners**, and the organizational unit they would belong to. The frontier labs hire these roles by the dozen. Salesforce hires them, as far as we can tell from public requisitions, not at all under those titles. The work that would prevent the failures the Forward Deployed Engineers are deployed to fix is not being staffed.

This is the org-chart diagnostic from Chapter 5, applied to Salesforce specifically. The evaluation function does not exist as an organizational unit. The hiring pattern confirms it. The remediation pattern confirms it. The release process confirms it. There is no missing detail; the structure is what it is.

The Builder Gap, recast

Our prior writing has documented the Builder Gap — the tension at TDX 2026 between low-code admin tooling and the developer experience required to build production-grade agents. The framing in that work was developer experience and tooling. We want to recast it briefly here, because the Builder Gap is, structurally, an eval-deficit problem in different vocabulary.

When admins can build agents but cannot validate that those agents work correctly, the result is the same as when developers ship code without tests: confidence runs ahead of capability. The Salesforce ecosystem has solved the build-it side of the agent problem reasonably well. It has

not solved the know-whether-it-works side. That is what evals do, and it is what is missing. Calling it a Builder Gap names the symptom; calling it an eval-culture gap names the cause.

The fix for the Builder Gap, accordingly, is not better builder tooling alone. It is builder tooling paired with customer-facing evaluation tooling. The first lets admins create agents quickly; the second tells them whether the agent they created works. Without both, the platform is generating agents at scale and shipping them to production without the discipline to know which ones are reliable. Some fraction will be reliable, some will be deeply unreliable, and the customer experience will be the variance — which is exactly what the field reports show.

The cost of inaction in Spring 2026

This is the part we say plainly because the rest of the chapter only matters if the consequences are clear.

The cost of not closing the eval-culture gap shows up in three places, on three timelines.

Renewal risk. Agentforce is roughly \$800 million in annual recurring revenue, with adoption sitting around 5.3% of the broader Salesforce customer base. Some material fraction of those customers — we estimate, based on the field reports, perhaps a third — are sitting in the bad half of the hallucination variance. They do not yet know they are. They are discovering it through customer complaints, through internal reviews, through quiet conversations with their account teams. A year from now, those are renewal risk. Two years from now, they are canceled contracts. The cost is not theoretical; it is a slow bleed that compounds because the customers who experience the failures tell other customers, and the procurement narrative starts to shift.

Brand exposure. The next high-profile Agentforce failure that hits LinkedIn with a recognizable customer name attached — the kind of failure where an agent confidently quotes a wrong policy to a customer who screenshots it — is a brand event. Without a pre-existing eval narrative to point to, the response is purely defensive. With a pre-existing eval narrative — published methodology, customer-facing tooling, transparent methodology — the response is *“here is the instrumentation that catches this; here is what we are improving.”* The latter is a controllable conversation. The former is not.

Competitive exposure. The first vendor to publish credible neutral evaluation data for CRM agents reframes the procurement conversation, regardless of who actually has the better product. Microsoft’s Copilot Studio, or a hyperscaler-native MCP implementation, or a well-funded startup, or a coalition of AI labs and enterprise platforms — any of these could move first. The advantage of moving first goes to whoever does it. The cost of moving second is producing comparable data under pressure, on someone else’s methodological terms, in a procurement narrative already framed by the first mover.

The historical analogy is the SOC2-by-default transition in enterprise SaaS, roughly 2014–2018. Companies that built security culture in advance lapped companies that bolted it on after a breach. The cost of bolting it on after the breach was, by published industry estimates, three to five times the cost of building it in. The brand damage during the transition cycle was, in many cases, permanent. Eval culture in Spring 2026 sits exactly where security culture sat in 2012: visible to insiders, invisible to executives, three years away from being a procurement-table requirement.

The diagnosis, in one paragraph

Salesforce has world-class research output and no eval culture. The CRM Benchmark dashboard is excellent and is not a product. The SCUBA benchmark is rigorous and does not gate releases. The Trust Layer is necessary and is not sufficient. The hiring posture is firefighting and not prevention. The competitive moat is the customer footprint and is, on the dimension of trustworthy AI, structurally exposed. None of this is unfixable. All of it is fixable on a deliberate schedule, starting in the next quarter, with a Phase 2 engagement of the kind Rosetta is designed to deliver. The cost of starting now is one engagement. The cost of starting after the first public failure is the brand.

The next chapter explains why Rosetta, specifically, is the program through which the closing of the gap should begin.



Chapter 7 — Why Rosetta Specifically Needs an EvalOps Layer

Chapter 6 diagnosed the gap. This chapter explains why Rosetta is the right program through which to begin closing it.

The argument has four moves. First, the System of Intent framing makes eval discipline *constitutive* of Rosetta's product, not optional. Second, Rosetta has structural advantages — the canonical inventory, the practitioner network, the producer-neutrality posture — that no other Salesforce program possesses, and that map directly onto the three lessons the eval economy has taught about durable competitive advantage. Third, the program architecture sequences eval work across multiple phases in a way that compounds rather than depletes. Fourth, the Phase 2 engagement specifically is sized and structured to be the right first step.

We will take these in order.

Rosetta as System of Intent: why eval is the product

In conventional enterprise software, evaluation is something you do *to* a product. The product is the application; the evaluation is a quality control activity adjacent to it. You can imagine the application existing without the evaluation; it would simply be a worse application.

Rosetta does not work this way. The Rosetta program has been described, in earlier Keenan Vision working papers and in conversation with Rahul Auradkar and Steve Fisher, as a System of Intent — a governed semantic layer that captures what enterprise users *mean* when they ask questions, not just what their applications return. The product *is* the calibrated mapping between practitioner intent and system response. Take away the calibration and you do not have a degraded version of Rosetta; you do not have Rosetta.

This is not a rhetorical flourish. It is a precise structural claim. A System of Intent is, by construction, an artifact of accumulated judgment about what good answers look like across thousands of scenarios. That accumulated judgment, captured in rubrics and verified through preference data, is the same artifact this book has been calling an eval. The two are the same object viewed from different angles. Rosetta-the-program produces Rosetta-the-evaluation-corpus as a byproduct of doing its core work, and the evaluation corpus is what makes Rosetta-the-product trustworthy enough to deploy.

The implication for Phase 2 is direct. Investing in EvalOps inside Rosetta is not bolting an evaluation function onto an existing platform. It is building the infrastructure that makes the platform itself coherent. The work is constitutive. The accounting consequence — that Phase 2's eval investment compounds into Rosetta's product moat — follows from the structural fact, not from a strategic narrative we are layering on top.

This distinguishes Rosetta from every other place inside Salesforce where eval discipline could be installed. Agentforce, Data 360, Service Cloud, Sales Cloud — for any of these, evaluation infrastructure would be additive: useful, expensive, and organizationally orphaned because no existing function owns it. For Rosetta, evaluation infrastructure is the program, with one of its components named as such.

The synthetic community of practice, in eval vocabulary

Our earlier writing has framed Rosetta as a synthetic community of practice — a structured way to capture the judgment of expert Salesforce architects across the ecosystem and make it available, governed and contextualized, to systems that need it. The framing has been useful for explaining the program's purpose to non-technical audiences. We want to recast it briefly here,

because the synthetic-community-of-practice idea is, structurally, an eval framework in different vocabulary.

Practitioners are evaluators. When an architect like Mehmet Orun reviews a proposed identity resolution approach and pronounces it sound, that is an evaluation. When Ian Gotts looks at a Data 360 schema design and says it will not survive the customer's actual data, that is an evaluation. The judgment these practitioners produce is structurally identical to the preference data Mercor pays \$300 per hour to acquire from generic SME pools. The difference is that Mehmet's and Ian's judgments are grounded in years of work on the specific systems Rosetta is built around, whereas Mercor's pool is grounded in general expertise that has to be calibrated for each new domain.

This is the structural advantage that no other enterprise platform can replicate. Salesforce, through the Rosetta program, has access to practitioner judgment that is impossible for a frontier lab or a generic data factory to acquire at any price. The Mercor playbook works because expert judgment is rare; the Rosetta version of the playbook works because *Salesforce-architect-tier* judgment is rarer, more domain-specific, and structurally embedded in a community Salesforce has spent twenty years cultivating.

The Phase 2 engagement is, in part, the first systematic test of this thesis. If 1,000 architect-tier preference pairs produced under the program's methodology yield inter-rater reliability and downstream training utility comparable to or better than what Mercor produces from generic SME pools, the thesis is validated and the program scales. The probability of that outcome, in our judgment, is high — but the validation is itself the right first step, because it converts a strategic narrative into a measurable claim.

The canonical inventory as ground truth

Eval work, to be useful, requires ground truth: a defensible answer to the question *what is the correct response in this scenario?* Most enterprise eval programs struggle to establish ground truth, because the underlying systems they are evaluating are complex enough that no single source of truth exists. The eval-economy vendors handle this by paying experts to *generate* ground truth from scratch, scenario by scenario. It works but it is expensive and lossy — each scenario is a one-off and the generated ground truth does not compound.

Rosetta has a different and rarer asset: the canonical inventory. The Cypher-first content pipeline, the governed semantic model on Neo4j, the formal entity-and-relationship structure that the program has been building since v1.3 — all of this is, in eval-program vocabulary, a corpus of pre-existing ground truth. When an evaluator scores a response about how Customer Profile relates to Unified Individual, they do not have to invent the ground truth from scratch.

They reference the canonical inventory, which already encodes the correct relationships, governance constraints, and provenance metadata.

This collapses the cost of eval generation by an order of magnitude relative to scenario-by-scenario approaches. It also makes the eval data unusually durable: ground truth that is encoded in a versioned graph survives model changes, retrieval changes, and rubric refinements in a way that scenario-specific ground truth does not. Every preference pair generated against the canonical inventory becomes a permanent data point about how the system performs against a stable reference.

The Phase 2 engagement is built around this pattern. Each of the 1,000 preference pairs uses the canonical inventory as the reference frame; each grading decision can be cross-checked against the inventory; each rubric criterion is calibrated against entities and relationships that already exist in the graph. The result is preference data that compounds across model generations and retrieval configurations — the kind of asset that justifies the cost of producing it across multiple use cycles, not just one.

Producer / governance / consumer separation

A structural feature of Rosetta’s architecture that has not been widely appreciated is its natural support for the producer-governance-consumer separation that mature eval programs require.

In a frontier lab, the team that *produces* eval data (the rubric authors, the graders, the calibration owners), the team that *governs* eval methodology (independence from product, release-gating authority, methodology defense), and the team that *consumes* eval results (the model trainers, the product organization, the deployment team) are kept organizationally distinct. The separation is what makes eval results credible — a result produced, governed, and consumed by the same team is not an evaluation; it is a self-assessment.

The Rosetta architecture, almost incidentally, supports this separation cleanly. The canonical inventory is the producer-side artifact. The governance layer — the schema constraints, the provenance tracking, the rubric stewardship that Phase 2 establishes — is the governance-side artifact. The consumer-side surfaces, including any eventual customer-facing dashboards or Agentforce integrations, are downstream of both. The architecture itself enforces the separation that other platforms have to install organizationally and frequently fail to maintain.

For Salesforce specifically, this matters because the producer-neutrality posture is the foundation for the leaderboard play we will describe shortly. If the producer-side artifacts (the canonical inventory, the rubrics, the verifier infrastructure) are governed by Keenan Vision under explicit independence covenants, and Salesforce sits in the consumer position rather than the producer position, the resulting evaluation data is credible to third parties in a way that no

Salesforce-produced benchmark currently is. The architecture supports the credibility; the institutional commitments operationalize it.

This is also why we have argued, in earlier conversations with Rahul and Steve, that Phase 2 should commit explicitly to a producer-neutrality covenant: the canonical inventory is governed such that practitioners outside Salesforce can contribute and benefit, with appropriate IP arrangements, without the artifact being construed as Salesforce-owned proprietary IP. The commitment is small at the start of Phase 2 and consequential by Phase 3.

The leaderboard play

The strategic implication of the producer-governance-consumer separation is the leaderboard play, and we want to name it directly because it is the largest single piece of long-term value the Rosetta program produces.

LMarena's \$1.7 billion Series A demonstrated that the first credible neutral evaluation infrastructure in a category captures a trust premium that is hard to dislodge. The CRM agent category, in Spring 2026, has no equivalent. There is no neutral platform where Agentforce, Microsoft Copilot Studio, ServiceNow's AI agents, and the hyperscaler-native MCP implementations are evaluated head-to-head on shared CRM tasks under shared methodology. The first organization to build that infrastructure, under credible neutrality covenants, captures the same trust premium for CRM that LMarena now commands for general LLMs.

Rosetta is structurally well-positioned to host this. The canonical inventory provides the shared task substrate. The architect community provides the practitioner-grader pool. Keenan Vision's independence from Salesforce — confirmed by the Foundation Sprint engagement structure rather than employment relationship — provides the neutrality. SalesforceDevops.net provides the publication channel. The pieces fit, and they do not fit anywhere else.

The leaderboard would not be a Phase 2 deliverable. It is a Phase 3 deliverable, contingent on the rubric IP and verifier infrastructure built in Phase 2 holding up under wider scrutiny. But the architectural choices made in Phase 2 either preserve the option for Phase 3 or foreclose it, and the foreclosure is irreversible. Building Phase 2 in a way that locks rubrics, verifiers, and calibration sets to Salesforce-only infrastructure would produce useful internal eval data and would close off the leaderboard play entirely. Building Phase 2 with the producer-neutrality covenant in place keeps the option open at no incremental cost.

The argument to Rahul and Steve — and the one we have been making in the Foundation Sprint sessions — is that the option preservation is itself worth the small additional structural commitment. Whether or not Salesforce eventually wants to participate in a CRM Arena, having

the option to host or co-host one is strategically valuable; foreclosing the option through narrow Phase 2 design would be a mistake that becomes visible only when a competitor moves first.

What Phase 2 specifically delivers

To translate the strategic argument into the concrete Phase 2 commitments, the engagement produces six durable artifacts, each of which has standalone value and each of which compounds across subsequent phases.

1. The 1,000 architect-tier preference pairs are the headline deliverable: graded preference data on D360 retrieval scenarios, with rubric-justified scoring, suitable for both evaluation of current Rosetta behavior and training signal for future improvement. This is what the engagement is named for and what is most legible to procurement.

2. The rubric IP is the more durable artifact. A versioned, calibrated, multi-criterion rubric for D360 scenarios, refined across the 1,000 grading decisions, becomes the measurement instrument for every future Rosetta release. This is the asset Surge AI is profitable on, in different domain.

3. The calibrated LLM-judge infrastructure converts the human-grading work into ongoing automated evaluation capability. Once the LLM judge has been calibrated against the human grader pool to inter-rater reliability above 0.7, the program can run a million additional evaluations against the same rubric for the cost of LLM compute. This is where the engagement's per-eval economics shift from \$300 to \$0.50.

4. The methodology document captures the program's eval methodology in a form publishable as a Keenan Vision working paper, co-authored with UC Berkeley if the academic partnership materializes. The publication serves three purposes: it establishes Salesforce's commitment to eval-driven development in a public forum, it positions Keenan Vision as the methodology owner for future engagements, and it creates citation infrastructure that compounds.

5. The inter-rater reliability analysis is the methodological proof point. Cohen's kappa across the architect grader pool, with sub-analyses by rubric criterion and scenario type, demonstrates that the program produces eval data of training-grade quality. This is the artifact procurement and academic reviewers will scrutinize most closely, and the one that — if it lands above 0.7 — converts the engagement from advisory to authoritative.

6. The Phase 3 readiness assessment is the transition document. It maps the rubric IP, verifier infrastructure, and grader pool onto the requirements for Phase 3 (production evals, customer-facing tooling, leaderboard infrastructure) and recommends the specific commitments

needed to make Phase 3 viable. Without this assessment, Phase 2 ends in success and Phase 3 starts from scratch; with it, Phase 3 begins with the assets already in place.

The point of listing the six artifacts is not procurement-document detail. It is to make visible that Phase 2 is not a one-shot eval engagement; it is the foundation layer of a multi-phase program in which each phase amortizes against the assets built in the previous one. Rahul and Steve should read the Phase 2 SOW with this in mind: the price tag is for the first phase of a four-phase program, not for an isolated deliverable.




Figure 7.1 — Rosetta architecture with the EvalOps layer in the governance position. Phase 2 installs the rubric IP, verifier infrastructure, grader pool, calibration set, and golden set; Phase 3 productizes the consumer-side tools and the neutrality-governed leaderboard.

The compounding case

The argument-from-compounding is the close. Most enterprise vendor engagements produce deliverables that depreciate on delivery. A consulting report is most valuable the day it lands; a one-time eval result decays as soon as the model under evaluation changes. The deliverables of Phase 2, by contrast, *appreciate* across the program lifecycle.

The rubric IP becomes more valuable as the canonical inventory grows, because more entities and relationships fall under its scope. The calibrated LLM judge becomes more valuable as model generations turn over, because the calibration carries forward and the cost of running each new model evaluation drops toward the cost of compute. The methodology document

becomes more valuable as it is cited, because citations compound credibility. The grader pool, formalized as the Keenan Vision Eval Bench, becomes more valuable as its members complete more engagements together, because their inter-rater reliability improves and the onboarding cost for future engagements drops.

The frontier labs figured this out. It is why they spend nine figures per year on eval data and consider it bargain pricing. The data is not the product; the methodology, the calibration, and the accumulated judgment around the data are the product, and they appreciate.

For Salesforce, the implication is that funding Phase 2 is funding the construction of an asset class the company does not currently possess. Mature eval infrastructure is a competitive moat the frontier labs have spent two years building. Salesforce does not have it. Rosetta is the program through which Salesforce can build it, on a deliberate schedule, starting with one Phase 2 engagement, sized to produce durable assets rather than one-time deliverables.

The next chapter is the Monday-morning version of this argument: the specific actions any leader can take in the next week, regardless of Phase 2 timing, that move the institutional response from awareness to commitment.



Chapter 8 — What to Do Monday Morning

We promised, in the introduction, that this book would land in a place where you could act on Monday morning. This is that chapter. It is short on purpose. The previous seven chapters did the persuasive and analytical work; this chapter does the operational work, and operational chapters are most useful when they fit on a single sheet of paper.

The recommendations below are sequenced. The first three are for any leader, regardless of Phase 2 timing or any specific engagement with Keenan Vision. The next group is for leaders considering the Phase 2 engagement directly. The closing section is the metrics commitment we believe Salesforce should make publicly by Q4 2026, which is the deadline imposed not by us but by the competitive dynamics this book has documented.

Three things to do this week

These three actions cost nothing, take less than two hours of executive time combined, and produce diagnostic information that determines everything that follows. Do not skip them in favor of more elaborate actions. Their value is precisely in their simplicity.

1. Convene a thirty-minute meeting with the Agentforce product team and ask one question. *What is our internal pass threshold on SCUBA, and who has authority to block a release based on it?* Notice carefully what happens. If an answer comes quickly with named numerical thresholds and a named individual, the gap is smaller than this book has suggested and you should congratulate the team. If the answer is vague, deflected, or requires follow-up, the gap is exactly what we have described and you have just diagnosed it without commissioning a single hour of consulting work. The diagnostic is the meeting itself.

2. Request the customer-facing version of the CRM Benchmark dashboard. Send the request through whichever channel normally produces internal product asks. Notice what comes back. The most likely response is a clarification — the dashboard exists internally, but no customer-facing version has been built. This is itself the answer. The Tableau workbook the AI Research team produces is a research artifact; the productized version is the artifact customers need, and its absence is a productization gap that has not been assigned an owner. The request you just made is the first step in assigning it.

3. Ask the procurement team how they would evaluate Salesforce's eval claims if they were buying Agentforce from a competitor. This is the most uncomfortable of the three because it inverts the perspective. Salesforce procurement evaluates competitors' claims; have they ever applied the same scrutiny to Salesforce's own? If the answer is no, you have surfaced a third structural feature: the same skepticism Salesforce procurement brings to vendor evaluation has not been turned inward, and the resulting confidence in internal eval claims is not earned in the way confidence in external vendors' claims is required to be earned.

The three actions together, in our experience advising large enterprise organizations, take a leader from a vague sense that something is missing to a precise diagnosis of what is missing in roughly two hours of meeting time. Do them this week. The rest of this chapter is harder to act on without that diagnostic groundwork.

The Phase 2 engagement, sized

For leaders considering whether to commission the Phase 2 engagement directly, the headline parameters are:

- **Scope:** 1,000 architect-tier preference pairs against the canonical D360 inventory, with the six durable artifacts named in Chapter 7 (preference pairs, rubric IP, calibrated LLM-judge, methodology document, inter-rater reliability analysis, Phase 3 readiness assessment).
- **Duration:** 120 calendar days, with mobilization in the first thirty and delivery across the remaining ninety.

- **Cost:** sized in the high six figures, structured to be comparable in magnitude to the existing Foundation Sprint engagement and to absorb the political and methodological friction that surfacing the gap will produce.
- **Co-funding option:** Salesforce supplying access to current D360 architects on payroll for the spot-check tier reduces loaded cost by roughly 30% and converts Salesforce from customer to co-investor, which is the right strategic posture for closing a gap of this kind.

The engagement is the first phase of a multi-phase program, not a one-shot deliverable. Read the Phase 2 SOW with that framing. Phase 3 contemplates production eval infrastructure and the leaderboard play; Phase 4 contemplates extension to Sales Cloud and Service Cloud. The price of Phase 2 is the foundation cost; the value compounds across the program lifecycle.

Five governance roles to create

Closing the eval-culture gap requires institutional roles that do not currently exist at Salesforce. We recommend creating five, in this order, over the next two quarters. Each can begin as a fractional commitment from existing leaders and grow into a dedicated role as the program matures.

Rubric Steward of Record. The named individual responsible for rubric authorship, versioning, and stewardship across the program. Single most important role; create it first. The Steward owns the rubric the way an editor owns a publication's voice.

Eval Lead. The leader of the evaluation function as an organizational unit, with the authority to commission evaluations and to surface results to executive leadership. Reports outside the product line; the org-chart point in Chapter 5 is the structural reason.

Calibration Owner. The individual responsible for inter-rater reliability, drift detection, and the quarterly recalibration cadence that keeps eval results comparable across model generations.

Ecosystem Steward. The individual responsible for the producer-neutrality covenants — managing the relationship between Salesforce and the practitioner community whose judgment populates the eval corpus, and ensuring the canonical inventory remains contributable from outside Salesforce.

Practitioner Council Chair. The convener of the standing architect council that informs rubric design, scenario selection, and methodology refinement. The role institutionalizes the synthetic community of practice as an operating function, not an aspirational one.

These five roles, named and staffed, are the organizational backbone of an eval culture. Without them, eval work is project-shaped and depletes. With them, eval work is program-shaped and compounds.

The Q4 2026 commitments

We close with a commitment Salesforce should make publicly, by Q4 2026, regardless of the engagement structure or vendor relationships involved. The commitment has three parts. We are direct about the deadline because the competitive dynamics described in Chapter 6 do not wait for internal readiness.

Pass-rate floor on a public benchmark. Salesforce should publicly commit to a minimum pass rate on a named benchmark — SCUBA is the natural choice — that all production Agentforce releases must clear. Anthropic’s Responsible Scaling Policy is the structural template. The number itself can be debated; the commitment to gate releases on it cannot be.

Hallucination ceiling on production traffic. Salesforce should commit to a maximum hallucination rate, measured by a defined methodology, that production deployments must remain below. Customers exceeding the ceiling — through configuration choices or grounding gaps — should receive instrumented feedback that allows them to remediate before customer-facing failures accumulate.

Calibration cadence published with results. Salesforce should commit to a quarterly publication cadence in which eval methodology, results, and rubric updates are released externally. The cadence is itself the commitment; the discipline of publishing on a schedule is what produces the institutional muscle that ad hoc disclosure does not.

The three commitments, taken together, are the move from a platform that does evaluation work to a platform that has installed eval culture. They are not exotic. They are within reach of any organization with the institutional will to make them. The question is whether Salesforce, in 2026, decides to be the platform that makes them — or whether a competitor moves first and reframes the procurement conversation around its own commitments.

Closing

We have written this book because Salesforce, in our judgment, has the talent, the customer footprint, and the strategic position to be the enterprise AI platform that defines what trustworthy looks like in this category. It does not yet have the cultural and organizational infrastructure to claim that position. The gap is real and it is closeable.

Eval culture is not a feature you ship. It is the structural muscle that determines whether everything else you ship can be trusted. Salesforce has the research talent. It has the customer footprint. It has not yet built the muscle. The cost of building it now is one engagement. The cost of building it after a public failure is the brand.

Phase 2 of the Rosetta program is the engagement. The Monday-morning actions in this chapter are how it begins. The decision to begin is yours.

We are at your service when you are ready.

— *Vernon Keenan, Keenan Vision · Spring 2026*



Glossary

Calibration set. A fixed reference set of responses with documented scores, used to train new graders to the same rubric standard and to detect rubric drift over time.

Capability eval. A structured measurement of whether a model can perform a representative task at all. Distinguished from reliability evals (consistency across edge cases) and preference evals (judgment between candidates).

Canonical inventory. In the Rosetta program, the authoritative graph-encoded set of entities, relationships, and provenance metadata that serves as ground truth for evaluation.

Cohen’s kappa (κ). A statistic measuring agreement between two graders, adjusted for the agreement expected by chance. Values: <0.40 poor, 0.40–0.60 moderate, 0.60–0.80 substantial, ≥ 0.80 near-perfect. The 0.7 floor is the conventional threshold for production-quality eval programs.

Constitutional AI. Anthropic’s methodology for training models against a written rubric (the “constitution”) that specifies criteria for desirable responses; a structural example of treating the rubric as the durable artifact and the model as the perishable one.

Contamination. When a benchmark’s test set has leaked into a model’s training data, invalidating the benchmark as a measure of capability.

Data factory. A vendor that produces expert-labeled training data and preference judgments at industrial scale. Examples: Mercor, Surge AI, Scale AI.

Drift, model. Changes in the system under evaluation that invalidate prior eval results.

Drift, rubric. Graders interpreting the same rubric differently as time passes, typically due to encountered edge cases.

Drift, domain. Changes in the underlying domain (customer language, regulations, product features) that invalidate the rubric or test scenarios.

Eval (evaluation). A structured measurement of model behavior against a defined rubric, with verifiable scoring, used to drive product decisions.

Eval-gated release. A release process in which models do not ship until they pass internal eval thresholds, with an organizational unit independent of the model-training team holding the authority to block release.

Eval arena. A platform that scores models against shared tasks under shared methodology, often with public leaderboards. Example: LMArena.

EvalOps. The operational discipline of running evaluation programs in production: rubric stewardship, calibration cadences, drift detection, customer-facing instrumentation.

Forward Deployed Engineer. A role focused on post-deployment troubleshooting of agent failures in customer environments. The hiring posture is firefighting; contrasted with eval engineering, which is preventive.

Frontier Safety Framework. Google DeepMind's published framework of capability levels, evaluation protocols, and deployment commitments contingent on evaluation outcomes.

Golden set. A versioned, contamination-resistant set of evaluation cases with documented correct answers, used to anchor program continuity across model generations.

Ground truth. A defensible answer to “what is the correct response in this scenario?” against which model outputs are scored.

Inter-rater reliability. The degree to which two or more graders, applying the same rubric to the same response, reach the same score. Measured by Cohen's kappa or related statistics.

LLM-as-judge. A verifier strategy in which a language model is prompted with a rubric and asked to score a response. Cost-effective at scale; requires calibration against human graders.

Preference eval. A structured measurement of which of two or more candidate responses an expert grader prefers, with rubric-justified scoring. The category that produces RLHF training data.

Preparedness Framework. OpenAI's published framework of capability thresholds and pre-defined safeguards triggered by evaluation outcomes.

Production eval. A continuous evaluation run against live user traffic, designed to detect drift, regression, and emergent failure modes that fixed test sets miss.

Rubric. A written specification of what counts as a correct or preferred response. The unit of disagreement in an eval program; the most expensive artifact to write well and the most durable when written well.

Reliability eval. A structured measurement of consistency across edge cases, designed to find failure modes rather than confirm peak performance.

Responsible Scaling Policy. Anthropic’s published policy committing the company to specific evaluation protocols at defined capability thresholds, with enforcement authority over deployment decisions.

RL gym. A sandboxed environment where AI agents practice multi-step tasks and earn deterministic reward signals. Examples: Mechanize, Prime Intellect.

RLHF (Reinforcement Learning from Human Feedback). The training methodology in which expert-graded preference data is used to refine model behavior on subjective tasks.

System of Intent. The Rosetta program’s framing as a governed semantic layer that captures what enterprise users mean when they ask questions.

Trust Layer. Salesforce’s downstream safety infrastructure (toxicity detection, PII redaction, prompt injection detection, grounding mechanisms). Necessary but distinct from upstream quality measurement; conflating the two is a category error.

Verifier. The mechanism that converts a response into a score. Three types: deterministic, LLM-as-judge, human expert.



References

Citations are organized by chapter. Where claims are based on primary public reporting, the source is named. Where claims rely on Foundation Sprint program documents or Keenan Vision research, the document title is referenced; readers with access to the program archive can locate the underlying material.

Chapter 1. Press, J. (2026). *On the structure of benchmarks*. Working paper. Salesforce AI Research, *SCUBA: Computer-Use Agents for Enterprise CRM Software* (October 2025), arXiv preprint.

Chapter 2. TechCrunch (June 2025), *Meta takes 49% stake in Scale AI in \$14.3B transaction*. Forbes (October 2025), *Mercor closes \$350M Series C at \$10B valuation*. Sacra equity research, *Mercor at \$500M ARR* (April 2026). The Information (December 2025), *Surge AI in talks to raise at \$25 billion*. SemiAnalysis (January 2026), *The RL Environment Economy*. Contrary Research, *LMarena business breakdown* (February 2026). NGP Capital (October 2025), *EvalOps: The New Differentiator*.

Chapter 3. General Reasoning (2026), *The OpenReward Standard*. Zhang et al. (2026), *EnterpriseOps-Gym: Multi-Tool Reasoning at Scale*. Landis, J. R. and Koch, G. G. (1977), *The Measurement of Observer Agreement for Categorical Data*, Biometrics 33(1).

Chapter 4. Salesforce AI Research, CRMArena and SCUBA technical reports (2024–2025). Mechanize and Prime Intellect public materials (2025–2026).

Chapter 5. Anthropic, *Responsible Scaling Policy* (multiple versions, 2023–2026). OpenAI, *Preparedness Framework* (2024). Google DeepMind, *Frontier Safety Framework* (2024). Bai et al. (2022), *Constitutional AI: Harmlessness from AI Feedback*. Public hiring requisitions across frontier labs and Salesforce, surveyed Spring 2026.

Chapter 6. Salesforce CRM Benchmark Tableau dashboard (2025–2026). Field reports on Agentforce hallucination rates aggregated from G2 and customer interviews. Salesforce job postings for Agentforce Forward Deployed Engineers, Spring 2026. Vernon Keenan, *The Builder Gap* (TDX 2026 coverage, SalesforceDevops.net).

Chapter 7. Vernon Keenan, *The Architect in the Graph* (Keenan Vision Working Paper #3, 2025). Vernon Keenan, *After the Quiet Erosion* (Keenan Vision memo, 2026). Foundation Sprint program documentation (Keenan Vision and Salesforce, 2025–2026, internal).

Chapter 8. This chapter draws on the analytical material established in Chapters 1–7; no additional external sources.



About Keenan Vision and the Rosetta Program

Keenan Vision LLC is an industry analysis and strategic advisory firm specializing in enterprise application development, the Salesforce ecosystem, and the economics of AI-driven labor transformation. The firm publishes SalesforceDevops.net, the leading independent publication for analysis of DevOps, AI, and enterprise technology trends in the Salesforce ecosystem, and produces research that has shaped industry conversations on Virtual Employee Economics, the Quiet Erosion thesis, and the architecture of agentic enterprise systems.

The Rosetta program is a Foundation Sprint engagement between Keenan Vision and Salesforce, building a governed semantic operating model — the Rosetta Enablement Agent — on Neo4j infrastructure. The program is sponsored on the Salesforce side by Rahul Auradkar and Steve Fisher and is positioned as the foundation layer for the next generation of enterprise AI deployment at Salesforce. Phase 2, scoped in this book, extends the program with an

EvalOps layer designed to install the cultural and organizational infrastructure that closes the gap diagnosed in Chapter 6.

For information about the Phase 2 engagement, the GTM Accelerator program for Salesforce ecosystem startups, or other Keenan Vision research and advisory services, contact the firm through SalesforceDevops.net.



About the Author

Vernon Keenan is CEO of Keenan Vision LLC and publisher of SalesforceDevops.net. He has tracked the enterprise software industry for over four decades, since the LISP era of artificial intelligence research, and serves as a strategic advisor to Marc Benioff and Salesforce executive leadership on AI strategy and competitive positioning. He coined the term Virtual Employee Economics and the Quiet Erosion thesis, and is the author of *Quantum Prometheus: The Kronos Singularity*. He is currently writing *Nexus of Abundance*, with a foreword committed by Marc Benioff. He lives in Oakland, California.



Acknowledgments

The author thanks Rahul Auradkar and Steve Fisher for their continued partnership in the Foundation Sprint engagement that has produced both the Rosetta program and the analytical foundation of this book. We thank Marc Benioff for his ongoing engagement with the Keenan Vision research program and for his commitment to the foreword of *Nexus of Abundance*.

We thank Prof. Abhishek Nagaraj of UC Berkeley Haas for the academic partnership that has informed the methodological rigor of our work on Virtual Employee Economics and, through that lens, the eval-economy framing in this book. We thank the broader Keenan Vision team — Chris Pearson, Alecia Wall, and Jon Robinson — for their work on the Rosetta platform that makes the Phase 2 EvalOps engagement viable.

We thank the Salesforce architect community whose practitioner judgment is the irreproducible asset this book is, in the end, about. Eval culture, in any organization, is built by individuals who insist on knowing whether the systems they build actually work. We are grateful to those

individuals and we have tried, in this book, to give the discipline they practice a vocabulary worthy of it.

— V.K., *Spring 2026*